

# Big Data 2

PROJEKAT 1

Veljko Veljović 1937

ELEKTRONSKI FAKULTET NIŠ

# Dataset - NYC Yellow Taxi Trip Data

- 1 Skup podataka o vožnjama žutih taksi vozila u New York-u
- 2 Podaci su prikupljeni od januara 2015. do marta 2016. godine.
- 3 Dataset je veličine 1.99 GB
- 4 Nije bilo potrebe za sređivanjem dataseta.

# producer.py

```
from confluent_kafka.admin import AdminClient, NewTopic
from confluent_kafka import Producer
from uuid import uuid4
import time
config = {
    "bootstrap.servers": 'localhost:9092'
}
topics_to_create = ["initial_data", "aggregated_results"]
new_topics = [NewTopic(topic, num_partitions=3, replication_factor=1) for topic in topics_to_create]

admin_client = AdminClient(config)

fs = admin_client.create_topics(new_topics)

producer = Producer(config)

file = open('yellow_tripdata_2015-01.csv', 'r', encoding='utf-8')

line = file.readline()
while line:
    producer.produce(
        topic="initial_data",
        key=str(uuid4()),
        value=line.encode("utf-8")
    )
    print(line)
    producer.flush()
    line = file.readline()
    time.sleep(3)

file.close()
```

# processor.py

Importi parsiranje argumenata i otvaranje spark sesije

```
from pyspark.sql import SparkSession
from pyspark.sql.types import StructType, StructField, IntegerType, DoubleType, StringType, TimestampType
from pyspark.sql.functions import from_csv, col, count, avg, stddev, window, sum, min, max, to_json, struct, lit, variance, expr, coalesce
import sys

window_duration = sys.argv[1]
slide_duration = sys.argv[2]
window_type = sys.argv[3]
filter_column = sys.argv[4]
filter_value = int(sys.argv[5])
group_column = sys.argv[6]
agg_column = sys.argv[7]
filter_condition = sys.argv[8]

spark = SparkSession.builder \
    .appName("Big-Data-2") \
    .config("spark.local.dir", "C:/spark_tmp") \
    .master("local[*]") \
    .getOrCreate()
spark.sparkContext.setLogLevel("ERROR")
```

šema podataka

```
taxi_schema = StructType([
    StructField("VendorID", IntegerType(), True),
    StructField("tpep_pickup_datetime", TimestampType(), True),
    StructField("tpep_dropoff_datetime", TimestampType(), True),
    StructField("passenger_count", IntegerType(), True),
    StructField("trip_distance", DoubleType(), True),
    StructField("pickup_longitude", DoubleType(), True),
    StructField("pickup_latitude", DoubleType(), True),
    StructField("RateCodeID", IntegerType(), True),
    StructField("store_and_fwd_flag", StringType(), True),
    StructField("dropoff_longitude", DoubleType(), True),
    StructField("dropoff_latitude", DoubleType(), True),
    StructField("payment_type", IntegerType(), True),
    StructField("fare_amount", DoubleType(), True),
    StructField("extra", DoubleType(), True),
    StructField("mta_tax", DoubleType(), True),
    StructField("tip_amount", DoubleType(), True),
    StructField("tolls_amount", DoubleType(), True),
    StructField("improvement_surcharge", DoubleType(), True),
    StructField("total_amount", DoubleType(), True)
])
schema_ddl = taxi_schema.simpleString()
```

# processor.py

## kreiranje dataframe-a

```
df = spark.readStream.format("kafka") \
    .option("kafka.bootstrap.servers", "localhost:9092") \
    .option("subscribe", "initial_data") \
    .option("failOnDataLoss", "false") \
    .load()

raw_df = df.selectExpr("CAST(value AS STRING) as csv_line")
parsed_df = raw_df.select(from_csv(col("csv_line"), schema_ddl).alias("data")).select("data.*")
```

## proračun

## definisanje uslova

```
if filter_condition.lower() == "greater":
    filtered_df = parsed_df.filter(col(filter_column) > filter_value)
elif filter_condition.lower() == "less":
    filtered_df = parsed_df.filter(col(filter_column) < filter_value)
else:
    filtered_df = parsed_df.filter(col(filter_column) == filter_value)
```

```
agg_df = filtered_df.groupBy(windowed_col, col(group_column)) \
    .agg(
        count("*").alias("count"),
        sum(agg_column).alias("sum"),
        min(agg_column).alias("min"),
        max(agg_column).alias("max"),
        avg(agg_column).alias("avg"),
        coalesce(stddev(agg_column), lit(0.0)).alias("stddev"),
        coalesce(variance(agg_column), lit(0.0)).alias("variance"),
        expr(f"percentile_approx({agg_column}, 0.5)").alias("median")
    )
agg_df = agg_df.withColumn(
    "condition",
    lit(f"Data where {filter_column} {filter_condition} {filter_value}")
)
```

# processor.py

## slanje podataka na novi topic i ispis na konzolu

```
non_empty_df = agg_df.filter(col("count") > 0)

standard_df = non_empty_df.withColumnRenamed(group_column, "group_value") \
| | | | | .withColumn("group_column_name", lit(group_column))

json_df = standard_df.select(to_json(struct("*")).alias("value"))

query_kafka = json_df.writeStream \
    .outputMode("complete") \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "localhost:9092") \
    .option("topic", "aggregated_results") \
    .option("checkpointLocation", "C:/spark_checkpoints/aggregated_results") \
    .option("failOnDataLoss", "false") \
    .trigger(processingTime="30 seconds") \
    .start()

query_console = standard_df.writeStream \
    .outputMode("complete") \
    .format("console") \
    .option("truncate", False) \
    .trigger(processingTime="30 seconds") \
    .start()

spark.streams.awaitAnyTermination()
```

# data-storage-service.py

## kreiranje keyspace i tabele

```
import json
from confluent_kafka import Consumer
from uuid import uuid4
from cassandra.cluster import Cluster

cluster = Cluster(['localhost'])
session = cluster.connect()

session.execute("""
    CREATE KEYSPACE IF NOT EXISTS taxi
        WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1}
""")
session.set_keyspace('taxi')

session.execute("""
    CREATE TABLE IF NOT EXISTS aggregated_taxi (
        window_start timestamp,
        window_end timestamp,
        group_column_name text,
        group_column_value int,
        condition text,
        count int,
        sum double,
        min double,
        max double,
        avg double,
        stddev double,
        variance double,
        median double,
        PRIMARY KEY ((window_start, window_end), min)
"""
)
```

## upis podataka u odgovarajuće tabele

```
def insert_aggregated_record(data, session):
    session.execute(
        """
        INSERT INTO aggregated_taxi (
            window_start, window_end, group_column_value, group_column_name,
            count, sum, min, max, avg, stddev, variance, median, condition
        ) VALUES (%s, %s, %s)
        """,
        (
            data.get("window", {}).get("start", None),
            data.get("window", {}).get("end", None),
            data.get("group_value", None),
            data.get("group_column_name", None),
            data.get("count", None),
            data.get("sum", None),
            data.get("min", None),
            data.get("max", None),
            data.get("avg", None),
            data.get("stddev", None),
            data.get("variance", None),
            data.get("median", None),
            data.get("condition", None)
        )
    )
```

# data-storage-service.py

## Subscribe na topike i insertovanje podataka u bazu

```
config = {
    "bootstrap.servers": 'localhost:9092',
    "group.id": f"consumer-{uuid4()}",
    "auto.offset.reset": "earliest"
}

consumer = Consumer(config)

consumer.subscribe(["aggregated_results"])
```

```
try:
    while True:
        msg = consumer.poll(1.0)
        if msg is None:
            continue
        if msg.error():
            print(f"Consumer error: {msg.error()}")
            continue

        message = msg.value().decode('utf-8')
        print(f"Received message: {message}")

        data = json.loads(message)
        insert_aggregated_record(data, session)

finally:
    consumer.close()
    print("Consumer closed")
```

# Testiranje

Filtrira taksi-vožnje sa manje od tri putnika i prebrojava ih, grupiše podatke u fiksne vremenske prozore od dva sata i izračunava statističke mere nad dužinom pređenog puta.

```
spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.3.1 processor.py "120 minutes" "60 minutes" sliding passenger_count 3 VendorID trip_distance less
```

```
C:\Users\veljk\OneDrive\Desktop\Big-Data-1>docker ps
CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              PORTS                 NAMES
b037305d2343        cassandra:latest   "docker-entrypoint.s..."   About a minute ago   Up     About a minute ago   0.0.0.0:9042->9042/tcp, [::]:9042->9042/tcp   cassandra-container
23e5bd8926c7        apache/kafka       "/__cacert_entrypoin..."   About a minute ago   Up     About a minute ago   0.0.0.0:9092->9092/tcp, [::]:9092->9092/tcp   kafka

C:\Users\veljk\OneDrive\Desktop\Big-Data-1>docker exec -it b037305d2343 bash
root@b037305d2343:/# cqlsh
Connected to Test Cluster at 127.0.0.1:9042
[cqlsh 6.2.0 | Cassandra 5.0.6 | CQL spec 3.4.7 | Native protocol v5]
Use HELP for help.
cqlsh> use taxi;
cqlsh:taxi> SELECT * FROM aggregated_taxi;

 window_start          | window_end            | min    | avg      | condition                                         | count | group_column_name | group_column_value | max    | median   | stddev   | sum     | va
riance
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
2015-01-15 18:00:00.000000+0000 | 2015-01-15 20:00:00.000000+0000 | 0.01 | 3.34263 | Data where passenger_count less 3 | 19 | VendorID | 2 | 18.06 | 2.37 | 4.01789 | 63.51 | 16
.14342
2015-01-15 18:00:00.000000+0000 | 2015-01-15 20:00:00.000000+0000 | 0.89 | 3.09636 | Data where passenger_count less 3 | 11 | VendorID | 2 | 7.13 | 2.38 | 2.19022 | 34.06 | 4
.79705
2015-01-15 18:00:00.000000+0000 | 2015-01-15 20:00:00.000000+0000 | 2.37 | 3.37167 | Data where passenger_count less 3 | 6 | VendorID | 2 | 7.13 | 2.38 | 1.90529 | 20.23 | 3
.63014
2015-01-15 12:00:00.000000+0000 | 2015-01-15 14:00:00.000000+0000 | 0.38 | 3.31941 | Data where passenger_count less 3 | 17 | VendorID | 2 | 15.2 | 1.51 | 4.67519 | 56.43 | 21
.85739
2015-01-15 12:00:00.000000+0000 | 2015-01-15 14:00:00.000000+0000 | 0.67 | 3.48571 | Data where passenger_count less 3 | 7 | VendorID | 2 | 15.2 | 1.53 | 5.21968 | 24.4 | 2
7.2451
2015-01-10 18:00:00.000000+0000 | 2015-01-10 20:00:00.000000+0000 | 0.3 | 2.83125 | Data where passenger_count less 3 | 16 | VendorID | 1 | 16.4 | 1.1 | 4.18278 | 45.3 | 17
.49562
2015-01-26 10:00:00.000000+0000 | 2015-01-26 12:00:00.000000+0000 | 0.5 | 1.54 | Data where passenger_count less 3 | 5 | VendorID | 1 | 4.3 | 1 | 1.55981 | 7.7 | 2.433
2015-01-04 12:00:00.000000+0000 | 2015-01-04 14:00:00.000000+0000 | 0.03 | 2.70667 | Data where passenger_count less 3 | 9 | VendorID | 2 | 8.98 | 1.26 | 3.53688 | 24.36 | 1
2.5095
2015-01-04 12:00:00.000000+0000 | 2015-01-04 14:00:00.000000+0000 | 0.5 | 1.26 | Data where passenger_count less 3 | 5 | VendorID | 1 | 2.5 | 1.1 | 0.792465 | 6.3 | 0.628
2015-01-04 12:00:00.000000+0000 | 2015-01-04 14:00:00.000000+0000 | 2.5 | 2.5 | Data where passenger_count less 3 | 1 | VendorID | 1 | 2.5 | 2.5 | 0 | 2.5 | 0
2015-01-24 22:00:00.000000+0000 | 2015-01-25 00:00:00.000000+0000 | 0.02 | 2.66053 | Data where passenger_count less 3 | 19 | VendorID | 2 | 10.2 | 2.22 | 2.29066 | 50.55 | 5
.24711

(11 rows)
cqlsh:taxi>
```

# Docker compose

```
services:  
kafka:  
  image: apache/kafka  
  container_name: kafka  
  ports :  
    - "9092:9092"  
  environment:  
    KAFKA_NODE_ID: 1  
    KAFKA_PROCESS_ROLES: broker,controller  
    KAFKA_LISTENERS: PLAINTEXT://0.0.0.0:29092,CONTROLLER://0.0.0.0:9093,PLAINTEXT_HOST://0.0.0.0:9092  
    KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://kafka:29092,PLAINTEXT_HOST://localhost:9092  
    KAFKA_CONTROLLER_LISTENER_NAMES: CONTROLLER  
    KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: CONTROLLER:PLAINTEXT,PLAINTEXT:PLAINTEXT,PLAINTEXT_HOST:PLAINTEXT  
    KAFKA_CONTROLLER_QUORUM_VOTERS: 1@localhost:9093  
    KAFKA_INTER_BROKER_LISTENER_NAME: PLAINTEXT  
    KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1  
    KAFKA_TRANSACTION_STATE_LOG_REPLICATION_FACTOR: 1  
    KAFKA_TRANSACTION_STATE_LOG_MIN_ISR: 1  
networks:  
  - bigdata-2  
  
cassandra:  
  image: cassandra:latest  
  container_name: cassandra-container  
  ports:  
    - "9042:9042"  
  environment:  
    - CASSANDRA_USER=admin  
    - CASSANDRA_PASSWORD=admin  
    - MAX_HEAP_SIZE=512M  
    - HEAP_NEWSIZE=100M  
  deploy:  
    resources:  
      limits:  
        memory: 2G  
  networks:  
    - bigdata-2
```

```
namenode:  
  image: bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8  
  container_name: namenode  
  restart: always  
  ports:  
    - "9870:9870"  
    - "9000:9000"  
  volumes:  
    - hadoop_namenode:/hadoop/dfs/name  
    - ./final_dataset:/data  
  environment:  
    - CLUSTER_NAME=test  
  env_file:  
    - ./hadoop.env  
  networks:  
    - bigdata-2  
  
datanode:  
  image: bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8  
  container_name: datanode  
  restart: always  
  volumes:  
    - hadoop_datanode:/hadoop/dfs/data  
  environment:  
    - SERVICE_PRECONDITION: "namenode:9870"  
  env_file:  
    - ./hadoop.env  
  networks:  
    - bigdata-2
```

# Docker compose

```
resourcemanager:  
  image: bde2020/hadoop-resourcemanager:2.0.0-hadoop3.2.1-java8  
  container_name: resourcemanager  
  restart: always  
  environment:  
    | SERVICE_PRECONDITION: "namenode:9000 namenode:9870 datanode:9864"  
  env_file:  
    | - ./hadoop.env  
  networks:  
    | - bigdata-2  
  
nodemanager:  
  image: bde2020/hadoop-nodemanager:2.0.0-hadoop3.2.1-java8  
  container_name: nodemanager  
  restart: always  
  environment:  
    | SERVICE_PRECONDITION: "namenode:9000 namenode:9870 datanode:9864 resourcemanager:8088"  
  env_file:  
    | - ./hadoop.env  
  networks:  
    | - bigdata-2  
  
historyserver:  
  image: bde2020/hadoop-historyserver:2.0.0-hadoop3.2.1-java8  
  container_name: historyserver  
  restart: always  
  environment:  
    | SERVICE_PRECONDITION: "namenode:9000 namenode:9870 datanode:9864 resourcemanager:8088"  
  volumes:  
    | - hadoop_historyserver:/hadoop/yarn/timeline  
  env_file:  
    | - ./hadoop.env  
  networks:  
    | - bigdata-2
```

```
spark-master:  
  image: bde2020/spark-master:3.1.2-hadoop3.2  
  container_name: spark-master  
  ports:  
    | - "8070:8080"  
    | - "7077:7077"  
  volumes:  
    | - ./spark-apps:/opt/spark-apps  
  environment:  
    | - INIT_DAEMON_STEP=setup_spark  
  networks:  
    | - bigdata-2  
  
spark-worker-1:  
  image: bde2020/spark-worker:3.1.2-hadoop3.2  
  container_name: spark-worker-1  
  depends_on:  
    | - spark-master  
  environment:  
    | - SPARK_MASTER=spark://spark-master:7077  
  networks:  
    | - bigdata-2  
  
spark-worker-2:  
  image: bde2020/spark-worker:3.1.2-hadoop3.2  
  container_name: spark-worker-2  
  depends_on:  
    | - spark-master  
  environment:  
    | - SPARK_MASTER=spark://spark-master:7077  
  networks:  
    | - bigdata-2  
  
volumes:  
  | hadoop_namenode:  
  | hadoop_datanode:  
  | hadoop_historyserver:  
  
networks:  
  | bigdata-2:  
    | driver: bridge
```

# Docker compose

## run.bat

```
@echo off
docker exec -i namenode bash -c "hdfs dfsadmin -safemode leave"
docker exec namenode hdfs dfs -rm -r /aggregated_results
docker exec spark-master /spark/bin/spark-submit --master spark://spark-master:7077 --packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.1.2 /opt/spark-apps/processor.docker.py
"120 minutes" "60 minutes" "tumbling" "passenger_count" "3" "VendorID" "trip_distance" "less"
pause
```

## docker app

```
spark = SparkSession.builder \
    .appName("Big-Data-2") \
    .getOrCreate()
```

```
df = spark.readStream.format("kafka") \
    .option("kafka.bootstrap.servers", "kafka:29092") \
    .option("subscribe", "initial_data") \
    .option("failOnDataLoss", "false") \
    .load()
```

```
query_kafka = json_df.writeStream \
    .outputMode("complete") \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "kafka:29092") \
    .option("topic", "aggregated_results") \
    .option("checkpointLocation", "hdfs://namenode:9000/aggregated_results") \
    .option("failOnDataLoss", "false") \
    .trigger(processingTime="30 seconds") \
    .start()
```

**Hvala na pažnji!**